

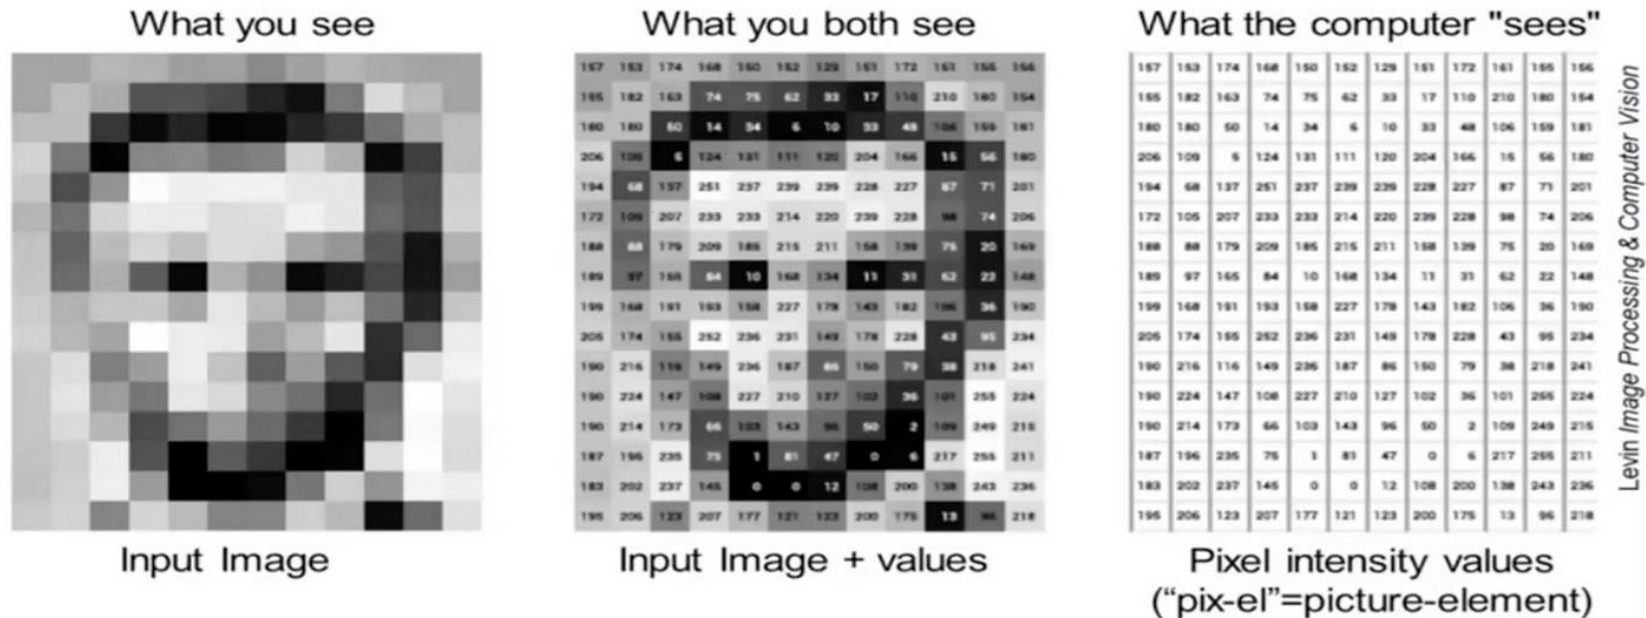
# Convolutional neural network (CNN)

---

HESAM HOSSEINI

SUMMER 2024

# How computer sees images



An image is just a matrix of numbers [0,255]!  
i.e., 1080x1080x3 for an RGB image

# How computer sees color

---



Red



Green




Blue




# How can we feed images into neural network

---



# A Problem

- Can a MLP identify two same types of flowers in these images?



# A Problem

---

We need a network that will activate, regardless of the exact location of the desired object?


We need shift invariance.



# Solution - Scanner

The entire operation can be viewed as a single giant network

- Composed of many “subnets” (one per window)
- With one key feature: all subnets are identical
- These are shared parameter networks



# Limitations of FNN

---

- For complex problems, we need multiple hidden layers in our FNN
  - Compounds the problem of having many weights
- Having too many weights
  - Makes learning more difficult as dimension of search space is increased
  - Makes training more time/resource consuming
  - Increases the likelihood of overfitting
- Problem is further compounded for color images
  - Each pixel in color image represented by 3 values (RGB color mode)
  - Since each pixel represented by 3 values, we say *channel* size is 3
  - Image represented by  $64 \times 64 \times 3 = 12,288$  values (rows x columns x channels)
  - Number of weights is now  $12,288 \times 500 = 6,144,000$





# Limitations of FNN

---

- Clear that FNN cannot scale to larger images (Too many weights)
- Another problem with FNN
  - 2D image represented as 1D vector in input layer
  - Any spatial relationship in the data is ignored


# From FCN to CNN

- Sparse Connection/interaction



# How to reduce number of parameters?

Parameter Sharing (plus Stride):




# Convolution/Correlation

- Linear Convolution (Linear Shift Invariant Systems)

$$y[n] = \sum_m x[m] h[n - m] = \sum_k h[k] x[n - k]$$

- Correlation:

$$y[n] = \sum_m x[m] h[n + m] = \sum_k h[k] x[n + k]$$




# Convolution

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved  
Feature



0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...	...	...	...	...	...	...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...	...	...	...	...	...	...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...	...	...	...	...	...	...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



308

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-498

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



164

+ 1 = -25


Bias = 1

Output

-25				...
				...
				...
				...
...	...	...	...	...

# Convolution shrinks


---






# padding

---







# Padding



Valid (Smaller)





Same (Equal)



Full (Larger)


# Padding

- a) Constant
- b) Replicate (nearest)
- c) Symmetric (mirror)
- d) Circular




# Stride


Convolution  
with Stride=1




Output



Convolution  
with Stride=2



Output



# Stride and padding


$$M = \frac{N + 2P - f}{stride} + 1$$

*Input* =  $N \times N$

*Padded Input* =  $(N + 2P) \times (N + 2P)$


*Filter* =  $f \times f$

*Output* =  $M \times M$



# Convolution Effect!

---




$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$





# Convolution Effect!

---




$$\begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$




# Convolution Effect!

---

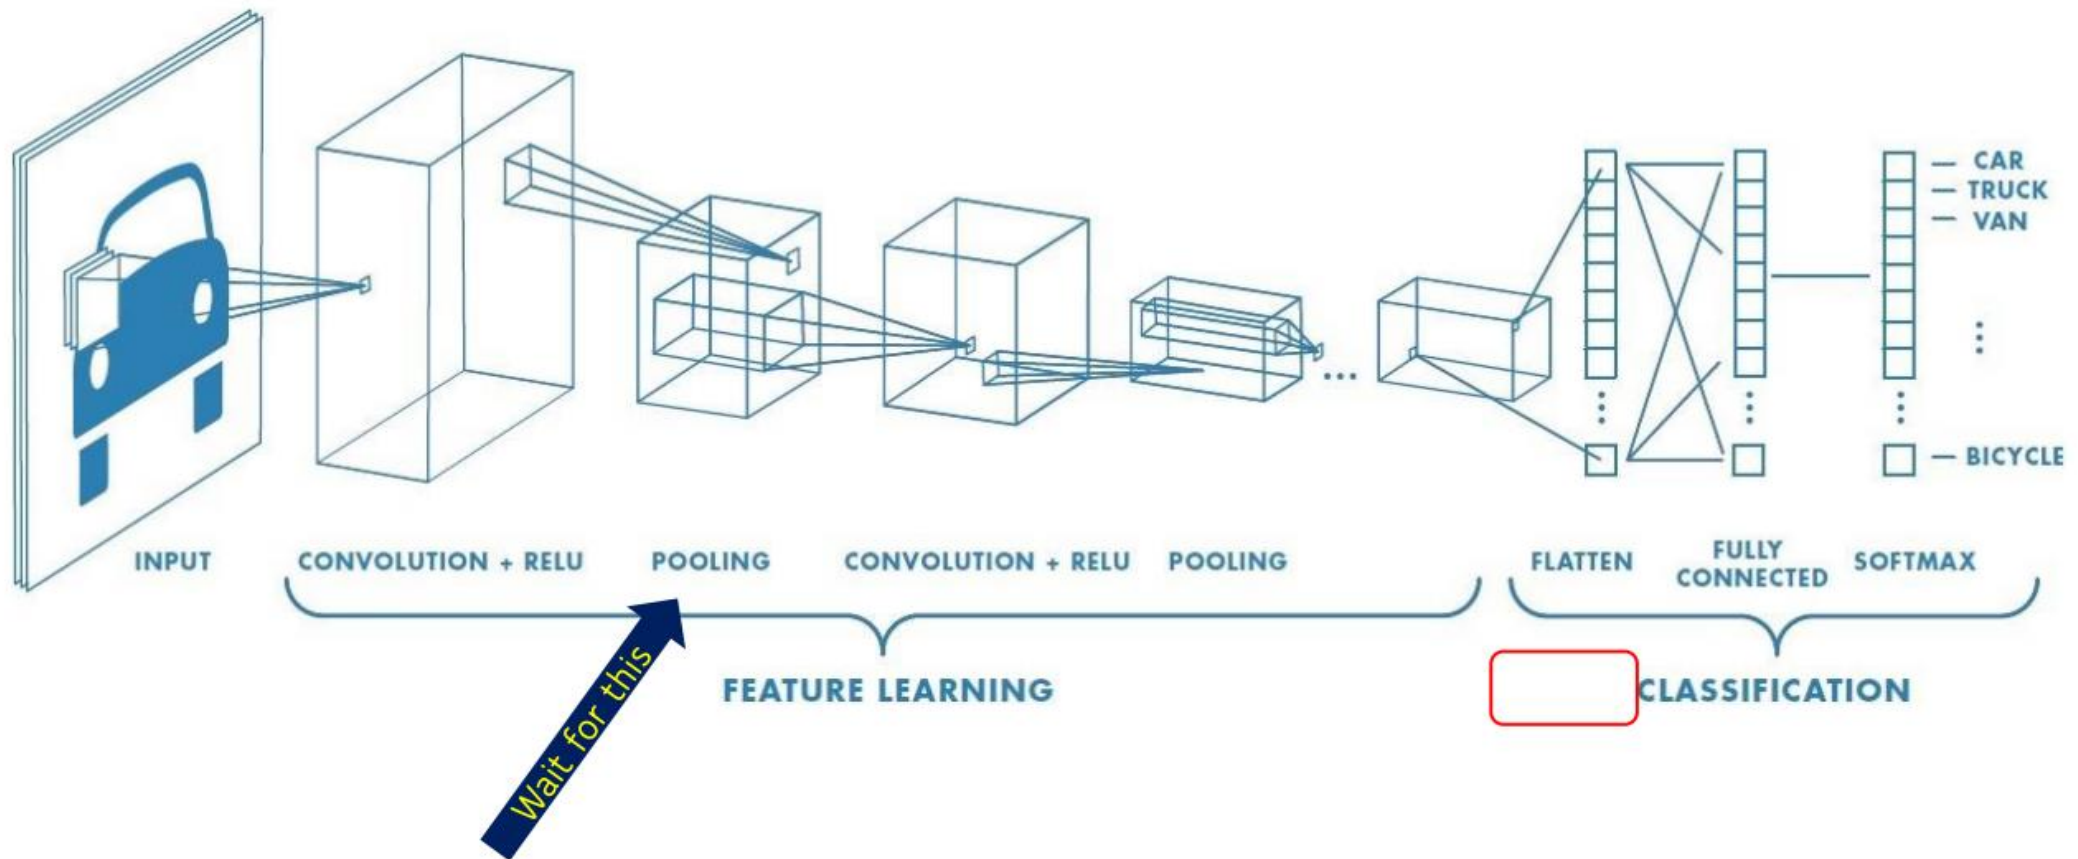


$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$





# A Typical CNN






# The Neuroscientific Basis

---

- In 1959 Hubel & Wiesel did an experiment to understand how visual cortex of brain processes visual info
  - Recorded activity of neurons in visual cortex of a cat
  - While moving a bright line in front of the cat
- Some cells fired when bright line is shown at a particular angle/location
  - Called these *simple* cells
- Other cells fired when bright line was shown regardless of angle/location
  - Seemed to detect movement
  - Called these *complex* cells
- Seemed complex cells receive inputs from multiple simple cells
  - Have an hierarchical structure
- Hubel and Wiesel won Noble prize in 1981

# The Neuroscientific Basis

Primary Visual Cortex, Theory: from 1959-1985




V1: Edge detection, etc.

V2: Extract simple visual properties (orientation, spatial frequency, color, etc)

V4: Detect object feature of intermediate complexity

IT: Object recognition

# CNN Visualization



Filter or weights



Baseball—or stripes?  
*mixed4a, Unit 6*



Animal faces—or snouts?  
*mixed4a, Unit 240*



Clouds—or fluffiness?  
*mixed4a, Unit 453*



Buildings—or sky?  
*mixed4a, Unit 492*




# Pooling

Modify and Subsample output of detector (ReLU):

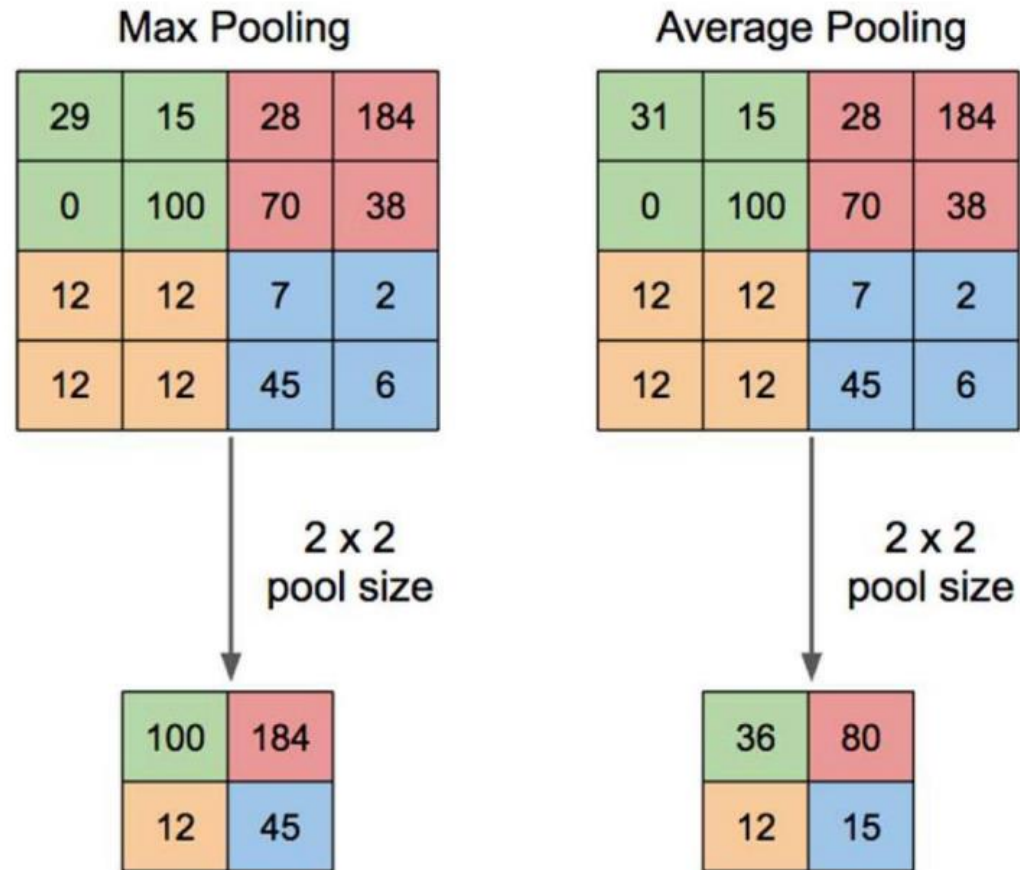
Using neighbors

- Max/min pooling
- Means pooling
- Median pooling

- Make Invariant to small variation (like shift)
- Reduce number of parameters
- Regularization



# Max Pooling vs Mean (Average) Pooling:





# A Challenging Problem: IMAGENET!

## Image Net

- 1000 Classes
- Train: 1.2 M
- test: 100 K

**airplane**



**automobile**



**bird**



**cat**




**deer**



**dog**




**frog**



**horse**




**ship**



**truck**



# performance on Image net



# AlexNet(2012)

---

## In Brief:

- ReLU (x6 faster than tanh)
- Regularization: 50% Dropout (x2 training time)
- 5 CL (Convolution Layer), 11×11, 5×5, 3×3
- 3 FC (Fully Connected), 4096-4096-1000
- MaxPool Layer: 3×3
- ReLu after all CL and FC
- Image Size: 256x256
- input Size: 227x227x3


## Details:

- First use of ReLU
- Normalize ReLu Output
- Several Data Augmentation
- Dropout: 0.5
- Batch Size: 128
- SGD Momentum: 0.9
- Learning rate 10<sup>-2</sup>
- 60M parameters




# AlexNet

Original (paper) Figure:



# AlexNet




# VGGnet

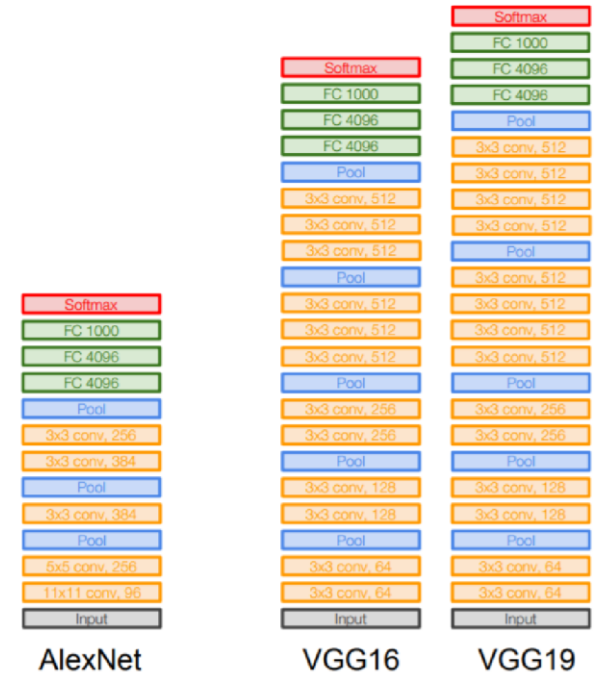
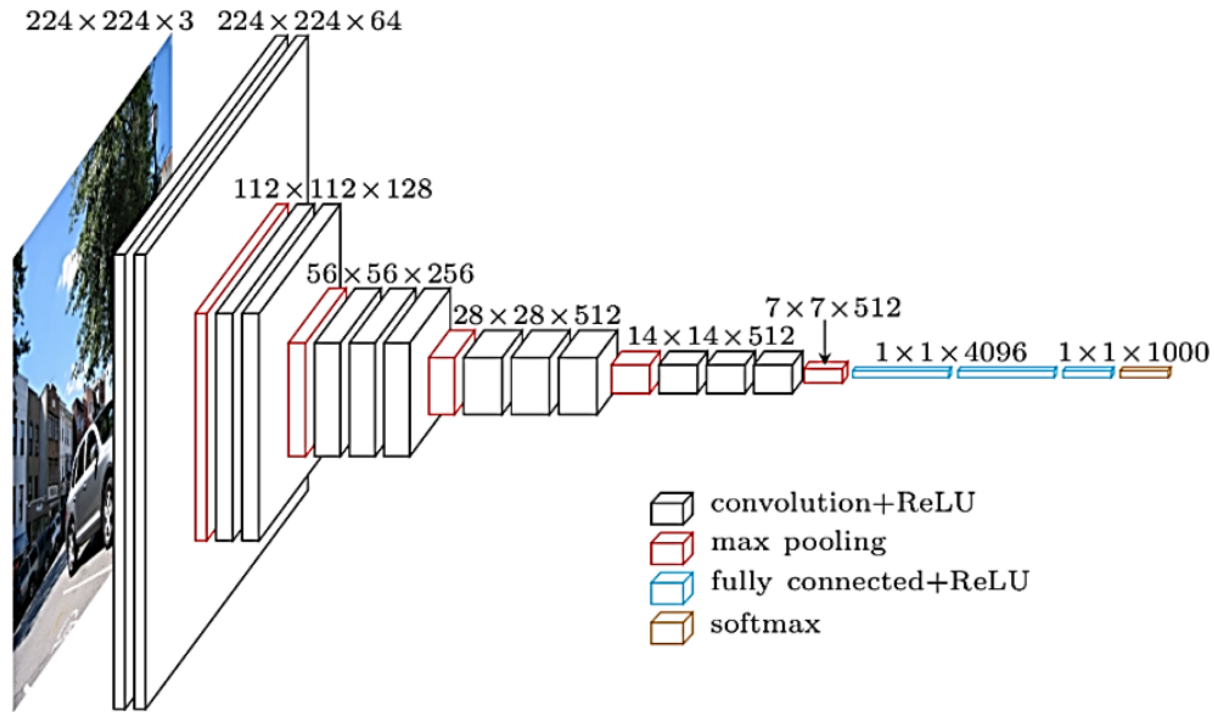
**VGG:** Visual Geometry Group

**Idea:** Smaller filter but deeper!

- 8 Layers in Alexnet → (16-19) ConvLayers
- ConvLayer (16/19): Only 3×3, Stride: 1, Padding: 1
  - More Non-linearity
- MaxPooling (5): 2×2, Stride: 2
- 138M Parameters



# VGGnet



# GoogLeNet

---

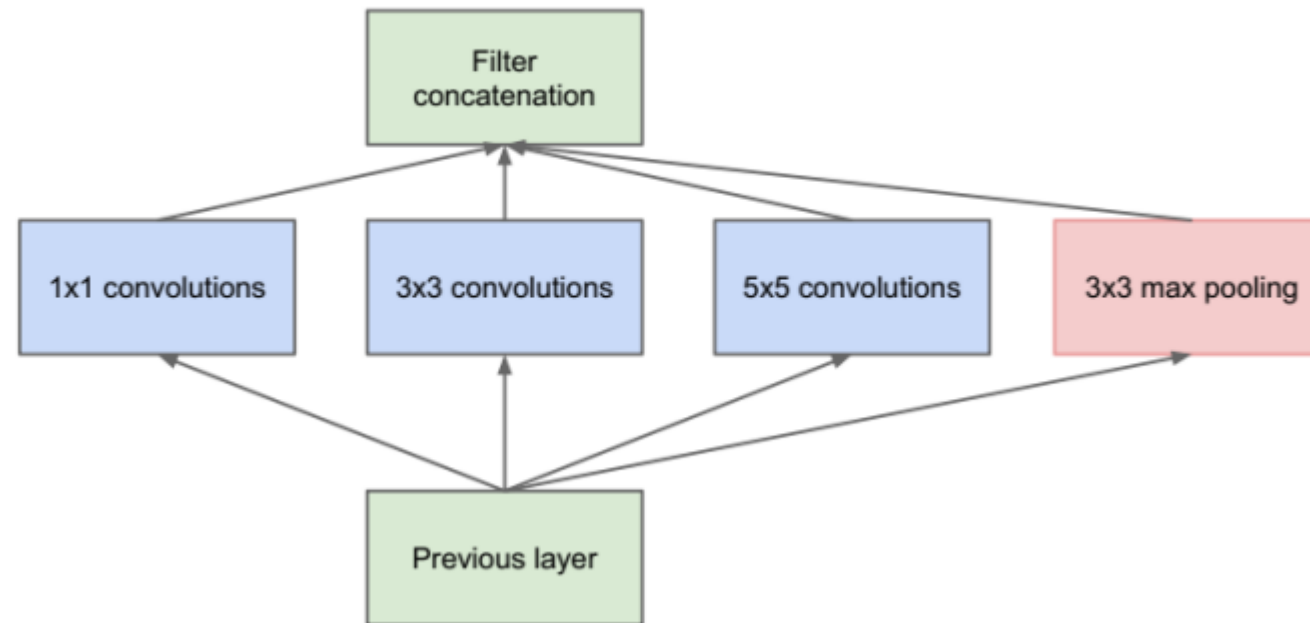
## Inception Module

### Idea and Motivation:

- Multiresolution Analysis (avoid cascade only convolution)
- Global Average Pooling before FC (dimension reduction)
- Bottleneck Layer (one dimension reduction)
- Auxiliary Classifier (shallow/weak output)


# GoogLeNet

Idea#1: Why Resolution reduction in successive Layer?



# GoogLeNet

Idea#2: Global Average Pooling Before FC




# GoogLeNet

## Idea #3: Auxiliary Classifier

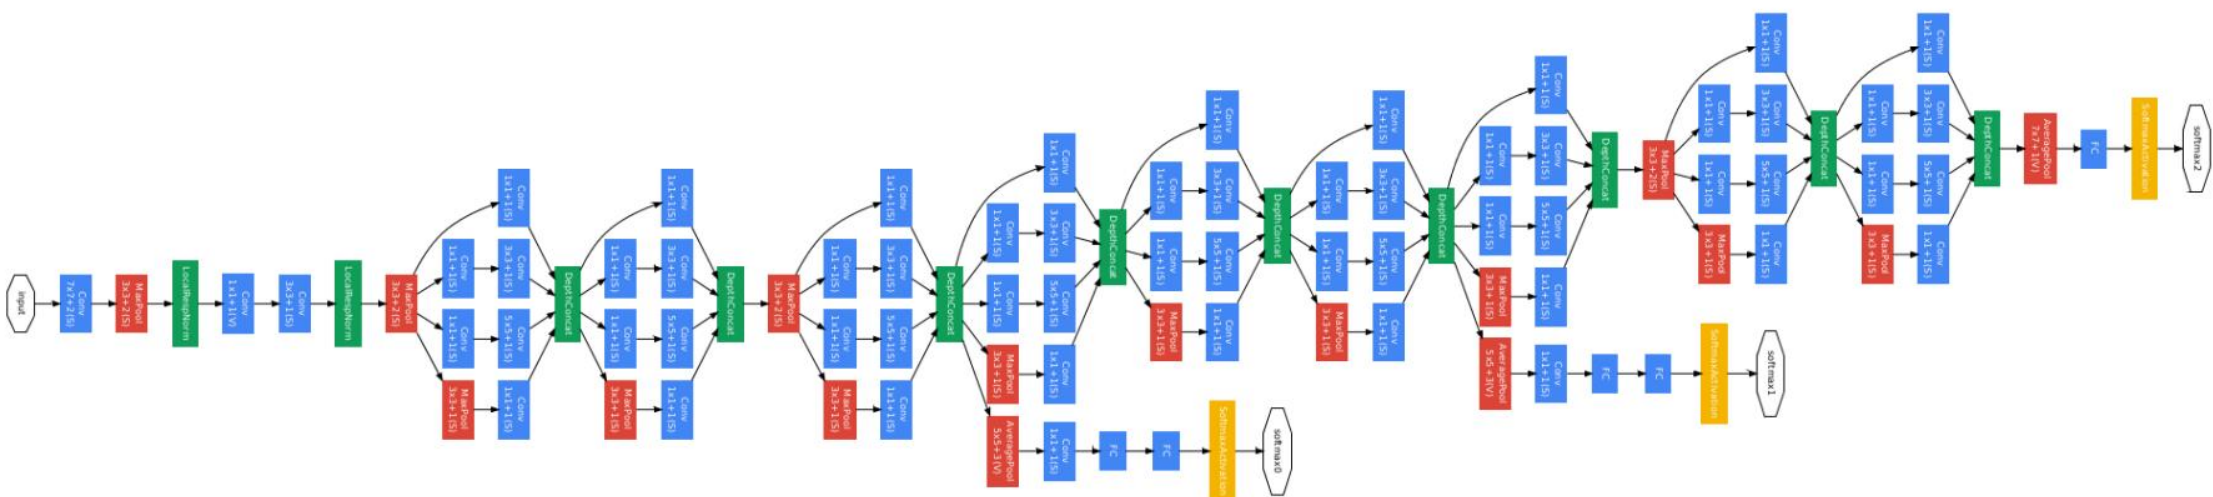
Too Deep (22 Layers) → Gradient flow problem!

Consider Shallow Outputs!







# GoogLeNet



# ResNet - Deep Residual Learning for Image Recognition (Microsoft)


## Problems with Deeper Networks:

- Performance degradation, Not due to overfitting
- Gradient vanishing/exploding



# ResNet

**Idea:** New Block (Shortcut Connection/Skip Connections)



# ResNet

---

A Identity map is hard to capture by highly nonlinear map.

In Residual Learning:

$$\mathbf{y} = \mathcal{F}(\mathbf{x}, \mathbf{W}) + \mathbf{x}$$

$H(x) = F(x) + x$ , the  $F(x)$  may learn to be zero!

- No Extra Parameter! 😊
- No Computational Complexity! 😊
- Problem of dimension  $F$  and  $X$ 
  - Solution:

$$\mathbf{y} = \mathcal{F}(\mathbf{x}, \mathbf{W}) + \mathbf{W}_s \mathbf{x}$$

# ResNet

---

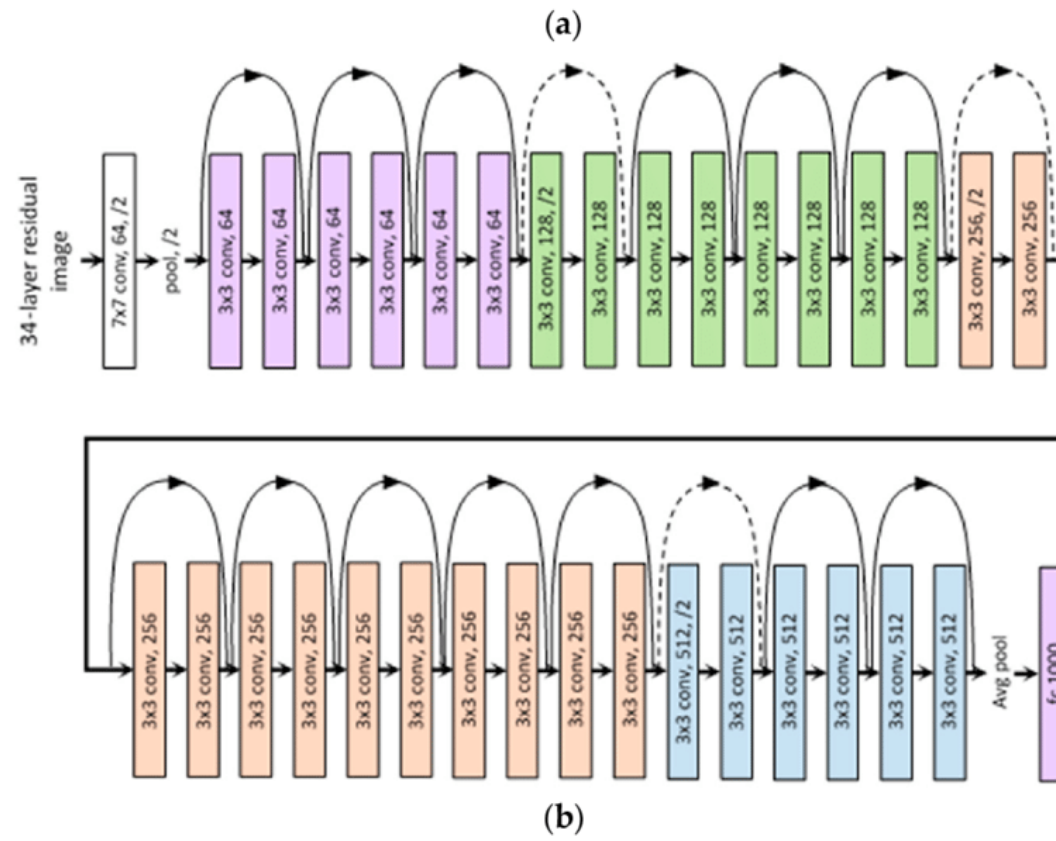
## ResNet Benefits:

- Vanishing and exploding gradients
- May increase # of layer


## Specification:

- Batch Normalization after each ConvLayer
- SGD + momentum (0.9)
- Minibatch Size: 256
- No Dropout

# ResNet



# DensNet




# Transfer Learning

## Traditional ML


vs

## Transfer Learning

- Isolated, single task learning:
  - Knowledge is not retained or accumulated. Learning is performed w.o. considering past learned knowledge in other tasks



- Learning of a new tasks relies on the previous learned tasks:
  - Learning process can be faster, more accurate and/or need less training data





# Transfer Learning

---

New dataset is **small** and **similar** to original dataset:

- Use final feature layer and train a linear or simple classifier

New dataset is **large** and **similar** to the original dataset:

- Fine-tune the full network

New dataset is **small** but very **different** from the original dataset:

- Train a SVM classifier from activations (middle layer, not final) in the network

New dataset is **large** and very **different** from the original dataset:

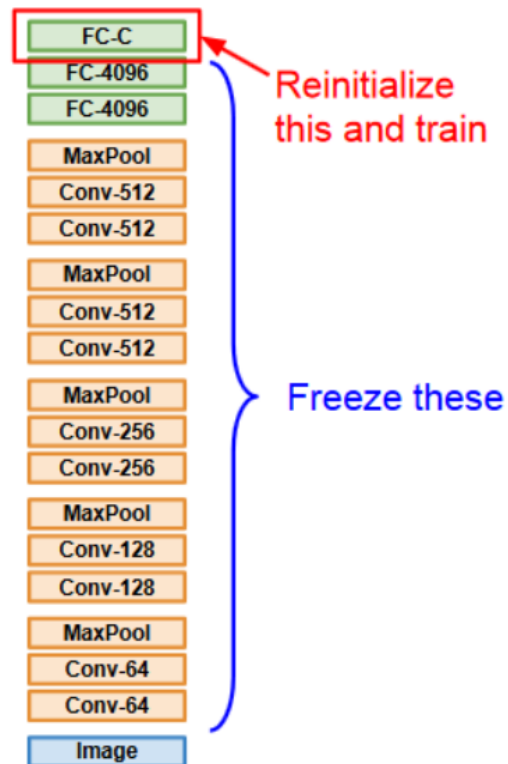
- Train from scratch (Initialize with weights from a pre-trained model)

# Transfer Learning

## 1. Train on Imagenet



## 2. Small Dataset (C classes)



## 3. Bigger dataset

